

Exploratory Data Analysis

Estimated time needed: **30** minutes

Objectives

After completing this lab you will be able to:

- Explore features or characteristics to predict price of car
- Analyze patterns and run descriptive statistical analysis
- Group data based on identified parameters and create pivot tables
- Identify the effect of independent attributes on price of cars

Import Data from Module 2

Import libraries:

```
#install specific version of libraries used in lab
#! mamba install pandas==1.3.3
#! mamba install numpy=1.21.2
#! mamba install scipy=1.7.1-y
#! mamba install seaborn=0.9.0-y

import pandas as pd
import numpy as np
import piplite
await piplite.install('seaborn')
```

<ipython-input-2-cec0a3d86e2d>:1: DeprecationWarning:
Pyarrow will become a required dependency of pandas in the next major
release of pandas (pandas 3.0),
(to allow more performant data types, such as the Arrow string type,
and better interoperability with other libraries)
but was not found to be installed on your system.
If this would cause problems for you,
please provide us feedback at
<https://github.com/pandas-dev/pandas/issues/54466>

```
import pandas as pd
```

Download the updated dataset by running the cell below.

The functions below will download the dataset into your browser and store it in dataframe `df`:

```
from pyodide.http import pyfetch

async def download(url, filename):
```

```

response = await pyfetch(url)
if response.status == 200:
    with open(filename, "wb") as f:
        f.write(await response.bytes())

file_path= "https://cf-courses-data.s3.us.cloud-object-
storage.appdomain.cloud/IBMDeveloperSkillsNetwork-DA0101EN-
SkillsNetwork/labs/Data%20files/automobileEDA.csv"

await download(file_path, "usedcars.csv")
file_name="usedcars.csv"

df = pd.read_csv(file_name, header=0)

```

Note: This version of the lab is working on JupyterLite, which requires the dataset to be downloaded to the interface. While working on the downloaded version of this notebook on their local machines(Jupyter Anaconda), the learners can simply **skip the steps above**, and simply use the URL directly in the `pandas.read_csv()` function. You can uncomment and run the statements in the cell below.

```

#filepath='https://cf-courses-data.s3.us.cloud-object-
storage.appdomain.cloud/IBMDeveloperSkillsNetwork-DA0101EN-
SkillsNetwork/labs/Data%20files/automobileEDA.csv'
#df = pd.read_csv(filepath, header=None)

```

View the first 5 values of the updated dataframe using `dataframe.head()`

```
df.head()
```

	symboling	normalized-losses	make	aspiration	num-of-
0	3	122	alfa-romero	std	two
1	3	122	alfa-romero	std	two
2	1	122	alfa-romero	std	two
3	2	164	audi	std	four
4	2	164	audi	std	four

	body-style	drive-wheels	engine-location	wheel-base	length	...
0	convertible	rwd	front	88.6	0.811148	...
1	convertible	rwd	front	88.6	0.811148	...
2	hatchback	rwd	front	94.5	0.822681	...

3	sedan	fwd	front	99.8	0.848630	...
4	sedan	4wd	front	99.4	0.848630	...

	compression-ratio	horsepower	peak-rpm	city-mpg	highway-mpg
price \					
0	9.0	111.0	5000.0	21	27
13495.0					
1	9.0	111.0	5000.0	21	27
16500.0					
2	9.0	154.0	5000.0	19	26
16500.0					
3	10.0	102.0	5500.0	24	30
13950.0					
4	8.0	115.0	5500.0	18	22
17450.0					

	city-L/100km	horsepower-binned	diesel	gas
0	11.190476	Medium	0	1
1	11.190476	Medium	0	1
2	12.368421	Medium	0	1
3	9.791667	Medium	0	1
4	13.055556	Medium	0	1

[5 rows x 29 columns]

Analyzing Individual Feature Patterns Using Visualization

To install Seaborn we use pip, the Python package manager.

Import visualization packages "Matplotlib" and "Seaborn". Don't forget about "%matplotlib inline" to plot in a Jupyter notebook.

```
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

# list the data types for each column
print(df.dtypes)
```

symboling	int64
normalized-losses	int64
make	object
aspiration	object
num-of-doors	object
body-style	object
drive-wheels	object
engine-location	object
wheel-base	float64

```

length          float64
width           float64
height          float64
curb-weight     int64
engine-type     object
num-of-cylinders object
engine-size     int64
fuel-system     object
bore            float64
stroke          float64
compression-ratio float64
horsepower      float64
peak-rpm        float64
city-mpg        int64
highway-mpg     int64
price           float64
city-L/100km    float64
horsepower-binned object
diesel          int64
gas             int64
dtype: object

```

Write your code below and press Shift+Enter to execute

```
df['peak-rpm'].dtypes
```

```
dtype('float64')
```

For example, we can calculate the correlation between variables of type "int64" or "float64" using the method "corr":

```

numeric_df = df.select_dtypes(include=['number'])
correlation_matrix = numeric_df.corr()
print(correlation_matrix)

```

	symboling	normalized-losses	wheel-base	length
\symboling	1.000000	0.466264	-0.535987	-0.365404
normalized-losses	0.466264	1.000000	-0.056661	0.019424
wheel-base	-0.535987	-0.056661	1.000000	0.876024
length	-0.365404	0.019424	0.876024	1.000000
width	-0.242423	0.086802	0.814507	0.857170
height	-0.550160	-0.373737	0.590742	0.492063
curb-weight	-0.233118	0.099404	0.782097	0.880665

engine-size	-0.110581	0.112360	0.572027	0.685025
bore	-0.140019	-0.029862	0.493244	0.608971
stroke	-0.008245	0.055563	0.158502	0.124139
compression-ratio	-0.182196	-0.114713	0.250313	0.159733
horsepower	0.075819	0.217299	0.371147	0.579821
peak-rpm	0.279740	0.239543	-0.360305	-0.285970
city-mpg	-0.035527	-0.225016	-0.470606	-0.665192
highway-mpg	0.036233	-0.181877	-0.543304	-0.698142
price	-0.082391	0.133999	0.584642	0.690628
city-L/100km	0.066171	0.238567	0.476153	0.657373
diesel	-0.196735	-0.101546	0.307237	0.211187
gas	0.196735	0.101546	-0.307237	-0.211187

	width	height	curb-weight	engine-size
bore \ symboling 0.140019	-0.242423	-0.550160	-0.233118	-0.110581 -
normalized-losses 0.029862	0.086802	-0.373737	0.099404	0.112360 -
wheel-base 0.493244	0.814507	0.590742	0.782097	0.572027
length 0.608971	0.857170	0.492063	0.880665	0.685025
width 0.544885	1.000000	0.306002	0.866201	0.729436
height 0.180449	0.306002	1.000000	0.307581	0.074694
curb-weight 0.644060	0.866201	0.307581	1.000000	0.849072
engine-size 0.572609	0.729436	0.074694	0.849072	1.000000
bore 1.000000	0.544885	0.180449	0.644060	0.572609
stroke 0.055390	0.188829	-0.062704	0.167562	0.209523 -
compression-ratio 0.001263	0.189867	0.259737	0.156433	0.028889
horsepower	0.615077	-0.087027	0.757976	0.822676

0.566936				
peak-rpm	-0.245800	-0.309974	-0.279361	-0.256733 -
0.267392				
city-mpg	-0.633531	-0.049800	-0.749543	-0.650546 -
0.582027				
highway-mpg	-0.680635	-0.104812	-0.794889	-0.679571 -
0.591309				
price	0.751265	0.135486	0.834415	0.872335
0.543155				
city-L/100km	0.673363	0.003811	0.785353	0.745059
0.554610				
diesel	0.244356	0.281578	0.221046	0.070779
0.054458				
gas	-0.244356	-0.281578	-0.221046	-0.070779 -
0.054458				

	stroke	compression-ratio	horsepower	peak-
rpm \				
symboling	-0.008245	-0.182196	0.075819	0.279740
normalized-losses	0.055563	-0.114713	0.217299	0.239543
wheel-base	0.158502	0.250313	0.371147	-0.360305
length	0.124139	0.159733	0.579821	-0.285970
width	0.188829	0.189867	0.615077	-0.245800
height	-0.062704	0.259737	-0.087027	-0.309974
curb-weight	0.167562	0.156433	0.757976	-0.279361
engine-size	0.209523	0.028889	0.822676	-0.256733
bore	-0.055390	0.001263	0.566936	-0.267392
stroke	1.000000	0.187923	0.098462	-0.065713
compression-ratio	0.187923	1.000000	-0.214514	-0.435780
horsepower	0.098462	-0.214514	1.000000	0.107885
peak-rpm	-0.065713	-0.435780	0.107885	1.000000
city-mpg	-0.034696	0.331425	-0.822214	-0.115413
highway-mpg	-0.035201	0.268465	-0.804575	-0.058598
price	0.082310	0.071107	0.809575	-0.101616
city-L/100km	0.037300	-0.299372	0.889488	0.115830

diesel	0.241303	0.985231	-0.169053	-0.475812
gas	-0.241303	-0.985231	0.169053	0.475812
	city-mpg	highway-mpg	price	city-L/100km
diesel \				
symboling	-0.035527	0.036233	-0.082391	0.066171 -
0.196735				
normalized-losses	-0.225016	-0.181877	0.133999	0.238567 -
0.101546				
wheel-base	-0.470606	-0.543304	0.584642	0.476153
0.307237				
length	-0.665192	-0.698142	0.690628	0.657373
0.211187				
width	-0.633531	-0.680635	0.751265	0.673363
0.244356				
height	-0.049800	-0.104812	0.135486	0.003811
0.281578				
curb-weight	-0.749543	-0.794889	0.834415	0.785353
0.221046				
engine-size	-0.650546	-0.679571	0.872335	0.745059
0.070779				
bore	-0.582027	-0.591309	0.543155	0.554610
0.054458				
stroke	-0.034696	-0.035201	0.082310	0.037300
0.241303				
compression-ratio	0.331425	0.268465	0.071107	-0.299372
0.985231				
horsepower	-0.822214	-0.804575	0.809575	0.889488 -
0.169053				
peak-rpm	-0.115413	-0.058598	-0.101616	0.115830 -
0.475812				
city-mpg	1.000000	0.972044	-0.686571	-0.949713
0.265676				
highway-mpg	0.972044	1.000000	-0.704692	-0.930028
0.198690				
price	-0.686571	-0.704692	1.000000	0.789898
0.110326				
city-L/100km	-0.949713	-0.930028	0.789898	1.000000 -
0.241282				
diesel	0.265676	0.198690	0.110326	-0.241282
1.000000				
gas	-0.265676	-0.198690	-0.110326	0.241282 -
1.000000				
	gas			
symboling	0.196735			
normalized-losses	0.101546			

wheel-base	-0.307237
length	-0.211187
width	-0.244356
height	-0.281578
curb-weight	-0.221046
engine-size	-0.070779
bore	-0.054458
stroke	-0.241303
compression-ratio	-0.985231
horsepower	0.169053
peak-rpm	0.475812
city-mpg	-0.265676
highway-mpg	-0.198690
price	-0.110326
city-L/100km	0.241282
diesel	-1.000000
gas	1.000000

The diagonal elements are always one; we will study correlation more precisely Pearson correlation in-depth at the end of the notebook.

```
df[['bore', 'stroke', 'compression-ratio', 'horsepower']].corr()
```

	bore	stroke	compression-ratio	horsepower
bore	1.000000	-0.055390	0.001263	0.566936
stroke	-0.055390	1.000000	0.187923	0.098462
compression-ratio	0.001263	0.187923	1.000000	-0.214514
horsepower	0.566936	0.098462	-0.214514	1.000000


Let's see several examples of different linear relationships:

Positive Linear Relationship

Let's find the scatterplot of "engine-size" and "price".

```
# Engine size as potential predictor variable of price
sns.regplot(x="engine-size", y="price", data=df)
plt.ylim(0,)
```

(0.0, 53529.821087193704)



We can examine the correlation between 'engine-size' and 'price' and see that it's approximately 0.87.


```
df[["engine-size", "price"]].corr()
```

	engine-size	price
engine-size	1.000000	0.872335
price	0.872335	1.000000

Highway mpg is a potential predictor variable of price. Let's find the scatterplot of "highway-mpg" and "price".

```
sns.regplot(x="highway-mpg", y="price", data=df)
```

```
<AxesSubplot:xlabel='highway-mpg', ylabel='price'>
```




We can examine the correlation between 'highway-mpg' and 'price' and see it's approximately -0.704.

```
df[['highway-mpg', 'price']].corr()
```

	highway-mpg	price
highway-mpg	1.000000	-0.704692
price	-0.704692	1.000000

Let's see if "peak-rpm" is a predictor variable of "price".

```
sns.regplot(x="peak-rpm", y="price", data=df)  
<AxesSubplot:xlabel='peak-rpm', ylabel='price'>
```



We can examine the correlation between 'peak-rpm' and 'price' and see it's approximately -0.101616.

```
df[['peak-rpm', 'price']].corr()
```

	peak-rpm	price
peak-rpm	1.000000	-0.101616
price	-0.101616	1.000000

Question 3 a):


Write your code below and press Shift+Enter to execute

```
df[["stroke", "price"]].corr()
```

	stroke	price
stroke	1.00000	0.08231
price	0.08231	1.00000


```
sns.regplot(x="stroke", y="price", data=df)
```

```
<AxesSubplot:xlabel='stroke', ylabel='price'>
```




Let's look at the relationship between "body-style" and "price".

```
sns.boxplot(x="body-style", y="price", data=df)  
<AxesSubplot:xlabel='body-style', ylabel='price'>
```




```
sns.boxplot(x="engine-location", y="price", data=df)  
<AxesSubplot:xlabel='engine-location', ylabel='price'>
```



Let's examine "drive-wheels" and "price".

```
# drive-wheels
sns.boxplot(x="drive-wheels", y="price", data=df)
<AxesSubplot:xlabel='drive-wheels', ylabel='price'>
```



Descriptive Statistical Analysis

This will show: the count of that variable the mean the standard deviation (std) the minimum value the IQR (Interquartile Range: 25%, 50% and 75%) the maximum value

We can apply the method "describe" as follows:

```
df.describe()
```

	symboling	normalized-losses	wheel-base	length
width \				
count	201.000000	201.000000	201.000000	201.000000
mean	0.840796	122.000000	98.797015	0.837102
std	1.254802	31.99625	6.066366	0.059213
min	-2.000000	65.000000	86.600000	0.678039
25%	0.000000	101.000000	94.500000	0.801538
50%	1.000000	122.000000	97.000000	0.832292
75%	2.000000	137.000000	102.400000	0.881788

```

0.925000
max      3.000000      256.00000  120.900000      1.000000
1.000000

      height  curb-weight  engine-size      bore      stroke  \
count  201.000000   201.000000   201.000000   201.000000   197.000000
mean    53.766667   2555.666667   126.875622     3.330692     3.256904
std      2.447822    517.296727    41.546834     0.268072     0.319256
min     47.800000   1488.000000    61.000000     2.540000     2.070000
25%     52.000000   2169.000000    98.000000     3.150000     3.110000
50%     54.100000   2414.000000   120.000000     3.310000     3.290000
75%     55.500000   2926.000000   141.000000     3.580000     3.410000
max     59.800000   4066.000000   326.000000     3.940000     4.170000

      compression-ratio  horsepower      peak-rpm      city-mpg
highway-mpg  \
count      201.000000   201.000000   201.000000   201.000000
201.000000
mean       10.164279   103.405534   5117.665368    25.179104
30.686567
std         4.004965    37.365700    478.113805     6.423220
6.815150
min         7.000000    48.000000   4150.000000    13.000000
16.000000
25%         8.600000    70.000000   4800.000000    19.000000
25.000000
50%         9.000000    95.000000   5125.369458    24.000000
30.000000
75%         9.400000   116.000000   5500.000000    30.000000
34.000000
max         23.000000   262.000000   6600.000000    49.000000
54.000000

      price  city-L/100km      diesel      gas
count  201.000000   201.000000   201.000000   201.000000
mean   13207.129353     9.944145     0.099502     0.900498
std    7947.066342     2.534599     0.300083     0.300083
min    5118.000000     4.795918     0.000000     0.000000
25%    7775.000000     7.833333     0.000000     1.000000
50%   10295.000000     9.791667     0.000000     1.000000
75%   16500.000000    12.368421     0.000000     1.000000
max   45400.000000    18.076923     1.000000     1.000000

```

The default setting of "describe" skips variables of type object. We can apply the method "describe" on the variables of type 'object' as follows:

```
df.describe(include=[ 'object' ])
```


	make	aspiration	num-of-doors	body-style	drive-wheels	\
count	201	201	201	201	201	
unique	22	2	2	5	3	
top	toyota	std	four	sedan	fwd	
freq	32	165	115	94	118	

	engine-location	engine-type	num-of-cylinders	fuel-system	\
count	201	201	201	201	
unique	2	6	7	8	
top	front	ohc	four	mpfi	
freq	198	145	157	92	

	horsepower-binned
count	200
unique	3
top	Low
freq	115

```
df['drive-wheels'].value_counts()
```

```
drive-wheels
fwd      118
rwd       75
4wd        8
Name: count, dtype: int64
```

We can convert the series to a dataframe as follows:

```
df['drive-wheels'].value_counts().to_frame()
```

	count
drive-wheels	
fwd	118
rwd	75
4wd	8

Let's repeat the above steps but save the results to the dataframe "drive_wheels_counts" and rename the column 'drive-wheels' to 'value_counts'.

```
drive_wheels_counts = df['drive-wheels'].value_counts().to_frame()
drive_wheels_counts.rename(columns={'drive-wheels': 'value_counts'},
inplace=True)
drive_wheels_counts
```

	count
drive-wheels	
fwd	118
rwd	75
4wd	8

Now let's rename the index to 'drive-wheels':

```
drive_wheels_counts.index.name = 'drive-wheels'
drive_wheels_counts
```

	count
drive-wheels	
fwd	118
rwd	75
4wd	8

We can repeat the above process for the variable 'engine-location'.

```
# engine-location as variable
engine_loc_counts = df['engine-location'].value_counts().to_frame()
engine_loc_counts.rename(columns={'engine-location': 'value_counts'},
inplace=True)
engine_loc_counts.index.name = 'engine-location'
engine_loc_counts.head(10)
```

	count
engine-location	
front	198
rear	3

Basics of Grouping

```
df['drive-wheels'].unique()
array(['rwd', 'fwd', '4wd'], dtype=object)
df_group_one = df[['drive-wheels', 'body-style', 'price']]
```

We can then calculate the average price for each of the different categories of data.

```
# Ensure all columns except 'drive-wheels' are numeric
numeric_columns = df_group_one.columns[df_group_one.dtypes !=
'object']
numeric_columns = numeric_columns.drop('drive-wheels')

# Convert these columns to numeric, ignoring errors
df_group_one[numeric_columns] =
df_group_one[numeric_columns].apply(pd.to_numeric, errors='coerce')

# Now perform the groupby operation
df_group_one = df_group_one.groupby(['drive-wheels'],
as_index=False).mean()
print(df_group_one)
```

```
-----  
-----  
KeyError                                Traceback (most recent call  
last)
```

```
Cell In[55], line 3
```

```
    1 # Ensure all columns except 'drive-wheels' are numeric  
    2 numeric_columns = df_group_one.columns[df_group_one.dtypes !=  
'object']  
----> 3 numeric_columns = numeric_columns.drop('drive-wheels')  
    5 # Convert these columns to numeric, ignoring errors  
    6 df_group_one[numeric_columns] =  
df_group_one[numeric_columns].apply(pd.to_numeric, errors='coerce')
```

```
File /lib/python3.12/site-packages/pandas/core/indexes/base.py:7069,  
in Index.drop(self, labels, errors)  
   7067 if mask.any():  
   7068     if errors != "ignore":  
-> 7069         raise KeyError(f"{labels[mask].tolist()} not found in  
axis")  
   7070     indexer = indexer[~mask]  
   7071 return self.delete(indexer)
```

```
KeyError: "['drive-wheels'] not found in axis"
```

```
# grouping results
```

```
df_gptest = df[['drive-wheels', 'body-style', 'price']]  
grouped_test1 = df_gptest.groupby(['drive-wheels', 'body-  
style'], as_index=False).mean()  
grouped_test1
```

	drive-wheels	body-style	price
0	4wd	hatchback	7603.000000
1	4wd	sedan	12647.333333
2	4wd	wagon	9095.750000
3	fwd	convertible	11595.000000
4	fwd	hardtop	8249.000000
5	fwd	hatchback	8396.387755
6	fwd	sedan	9811.800000
7	fwd	wagon	9997.333333
8	rwd	convertible	23949.600000
9	rwd	hardtop	24202.714286
10	rwd	hatchback	14337.777778
11	rwd	sedan	21711.833333
12	rwd	wagon	16994.222222

```
grouped_pivot = grouped_test1.pivot(index='drive-  
wheels', columns='body-style')  
grouped_pivot
```

```

      price
body-style convertible      hardtop      hatchback      sedan \
drive-wheels
4wd          NaN          NaN      7603.000000      12647.333333
fwd          11595.0      8249.000000      8396.387755      9811.800000
rwd          23949.6      24202.714286      14337.777778      21711.833333

body-style      wagon
drive-wheels
4wd          9095.750000
fwd          9997.333333
rwd          16994.222222

grouped_pivot = grouped_pivot.fillna(0) #fill missing values with 0
grouped_pivot

      price
body-style convertible      hardtop      hatchback      sedan \
drive-wheels
4wd          0.0          0.000000      7603.000000      12647.333333
fwd          11595.0      8249.000000      8396.387755      9811.800000
rwd          23949.6      24202.714286      14337.777778      21711.833333

body-style      wagon
drive-wheels
4wd          9095.750000
fwd          9997.333333
rwd          16994.222222

df_gptest2 = df[['body-style', 'price']]
grouped_test_bodystyle = df_gptest2.groupby(['body-style'], as_index=
False).mean()
grouped_test_bodystyle

   body-style      price
0  convertible  21890.500000
1    hardtop    22208.500000
2  hatchback    9957.441176
3    sedan     14459.755319
4    wagon     12371.960000

```

If you did not import "pyplot", let's do it again.

```

import matplotlib.pyplot as plt
%matplotlib inline


```

Let's use a heat map to visualize the relationship between Body Style vs Price.

```

#use the grouped results
plt.pcolor(grouped_pivot, cmap='RdBu')
plt.colorbar()
plt.show()

```



```

fig, ax = plt.subplots()
im = ax.pcolor(grouped_pivot, cmap='RdBu')

#label names
row_labels = grouped_pivot.columns.levels[1]
col_labels = grouped_pivot.index


#move ticks and labels to the center
ax.set_xticks(np.arange(grouped_pivot.shape[1]) + 0.5, minor=False)
ax.set_yticks(np.arange(grouped_pivot.shape[0]) + 0.5, minor=False)

#insert labels
ax.set_xticklabels(row_labels, minor=False)
ax.set_yticklabels(col_labels, minor=False)

#rotate label if too long
plt.xticks(rotation=90)

fig.colorbar(im)
plt.show()

```



Correlation and Causation

```
df.corr()
```

```
-----
-----
ValueError                                Traceback (most recent call
last)
```

```
Cell In[42], line 1
```

```
----> 1 df.corr()
```

```
File /lib/python3.12/site-packages/pandas/core/frame.py:11022, in
DataFrame.corr(self, method, min_periods, numeric_only)
```

```
    11020 cols = data.columns
```

```
    11021 idx = cols.copy()
```

```
> 11022 mat = data.to_numpy(dtype=float, na_value=np.nan, copy=False)
```

```
    11024 if method == "pearson":
```

```
    11025     correl = libalgos.nancorr(mat, minp=min_periods)
```

```
File /lib/python3.12/site-packages/pandas/core/frame.py:1981, in
DataFrame.to_numpy(self, dtype, copy, na_value)
```

```
    1979 if dtype is not None:
```

```

1980     dtype = np.dtype(dtype)
-> 1981 result = self._mgr.as_array(dtype=dtype, copy=copy,
na_value=na_value)
1982 if result.dtype is not dtype:
1983     result = np.array(result, dtype=dtype, copy=False)

File
/lib/python3.12/site-packages/pandas/core/internals/managers.py:1693,
in BlockManager.as_array(self, dtype, copy, na_value)
1691         arr.flags.writeable = False
1692     else:
-> 1693         arr = self._interleave(dtype=dtype, na_value=na_value)
1694         # The underlying data was copied within _interleave, so no
need
1695         # to further copy if copy=True or setting na_value
1697 if na_value is lib.no_default:

File
/lib/python3.12/site-packages/pandas/core/internals/managers.py:1752,
in BlockManager._interleave(self, dtype, na_value)
1750         else:
1751             arr = blk.get_values(dtype)
-> 1752             result[rl.indexer] = arr
1753             itemmask[rl.indexer] = 1
1755 if not itemmask.all():

ValueError: could not convert string to float: 'alfa-romero'

```

Sometimes we would like to know the significant of the correlation estimate.

P-value What is this P-value? The P-value is the probability value that the correlation between these two variables is statistically significant. Normally, we choose a significance level of 0.05, which means that we are 95% confident that the correlation between the variables is significant.

By convention, when the p-value is ≤ 0.001 : we say there is strong evidence that the correlation is significant. the p-value is ≤ 0.05 : there is moderate evidence that the correlation is significant. the p-value is ≤ 0.1 : there is weak evidence that the correlation is significant. the p-value is ≤ 0.1 : there is no evidence that the correlation is significant.

We can obtain this information using "stats" module in the "scipy" library.

```
from scipy import stats
```

Let's calculate the Pearson Correlation Coefficient and P-value of 'wheel-base' and 'price'.

```

pearson_coef, p_value = stats.pearsonr(df['wheel-base'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a
P-value of P =", p_value)

```

The Pearson Correlation Coefficient is 0.5846418222655085 with a P-value of $P = 8.076488270732338e-20$

Let's calculate the Pearson Correlation Coefficient and P-value of 'horsepower' and 'price'.

```
pearson_coef, p_value = stats.pearsonr(df['horsepower'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P = ", p_value)
```

The Pearson Correlation Coefficient is 0.8095745670036559 with a P-value of $P = 6.36905742825956e-48$

Let's calculate the Pearson Correlation Coefficient and P-value of 'length' and 'price'.

```
pearson_coef, p_value = stats.pearsonr(df['length'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P = ", p_value)
```

The Pearson Correlation Coefficient is 0.6906283804483643 with a P-value of $P = 8.016477466158871e-30$

Let's calculate the Pearson Correlation Coefficient and P-value of 'width' and 'price':

```
pearson_coef, p_value = stats.pearsonr(df['width'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_value )
```

The Pearson Correlation Coefficient is 0.7512653440522663 with a P-value of $P = 9.200335510485071e-38$

Conclusion:

Since the p-value is < 0.001 , the correlation between width and price is statistically significant, and the linear relationship is quite strong (~ 0.751).

Curb-Weight vs. Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'curb-weight' and 'price':

```
pearson_coef, p_value = stats.pearsonr(df['curb-weight'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P = ", p_value)
```

The Pearson Correlation Coefficient is 0.8344145257702845 with a P-value of $P = 2.1895772388939654e-53$

Let's calculate the Pearson Correlation Coefficient and P-value of 'engine-size' and 'price':


```
pearson_coef, p_value = stats.pearsonr(df['engine-size'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a
P-value of P =", p_value)
```

The Pearson Correlation Coefficient is 0.8723351674455188 with a P-value of P = 9.26549162219582e-64

Let's calculate the Pearson Correlation Coefficient and P-value of 'bore' and 'price':

```
pearson_coef, p_value = stats.pearsonr(df['bore'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a
P-value of P = ", p_value )
```

The Pearson Correlation Coefficient is 0.5431553832626601 with a P-value of P = 8.049189483935384e-17

We can relate the process for each 'city-mpg' and 'highway-mpg':

```
pearson_coef, p_value = stats.pearsonr(df['city-mpg'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a
P-value of P = ", p_value)
```

The Pearson Correlation Coefficient is -0.6865710067844684 with a P-value of P = 2.3211320655672357e-29

```
pearson_coef, p_value = stats.pearsonr(df['highway-mpg'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with
a P-value of P = ", p_value )
```

The Pearson Correlation Coefficient is -0.7046922650589532 with a P-value of P = 1.7495471144475574e-31

Conclusion:

Since the p-value is < 0.001, the correlation between highway-mpg and price is statistically significant, and the coefficient of about -0.705 shows that the relationship is negative and moderately strong.

Continuous numerical variables: Length Width Curb-weight Engine-size Horsepower City-mpg Highway-mpg Wheel-base Bore

Categorical variables: Drive-wheels

Thank you for completing this lab!

Author

Joseph Santarcangelo

Other Contributors

Mahdi Noorian PhD

Bahare Talayian

Eric Xiao

Steven Dong

Parizad

Hima Vasudevan

Fiorella Wener

Yi Yao.

Abhishek Gagneja

© IBM Corporation 2023. All rights reserved.

<!--

Change Log

| Date (YYYY-MM-DD) | Version | Changed By | Change Description |
|-------------------|---------|------------------|------------------------------------|
| 2023-09-28 | 2.2 | Abhishek Gagneja | Updated instructions |
| 2020-10-30 | 2.1 | Lakshmi | changed URL of csv |
| 2020-08-27 | 2.0 | Lavanya | Moved lab to course repo in GitLab |

--!>